

Lecture 13

# Rust and WebAssembly

# WebAssembly Demo

# What is wrong with JavaScript?

- The strength is a WEAKNESS:

**Easy to learn + Executes in browser + Dynamic Typing**

=

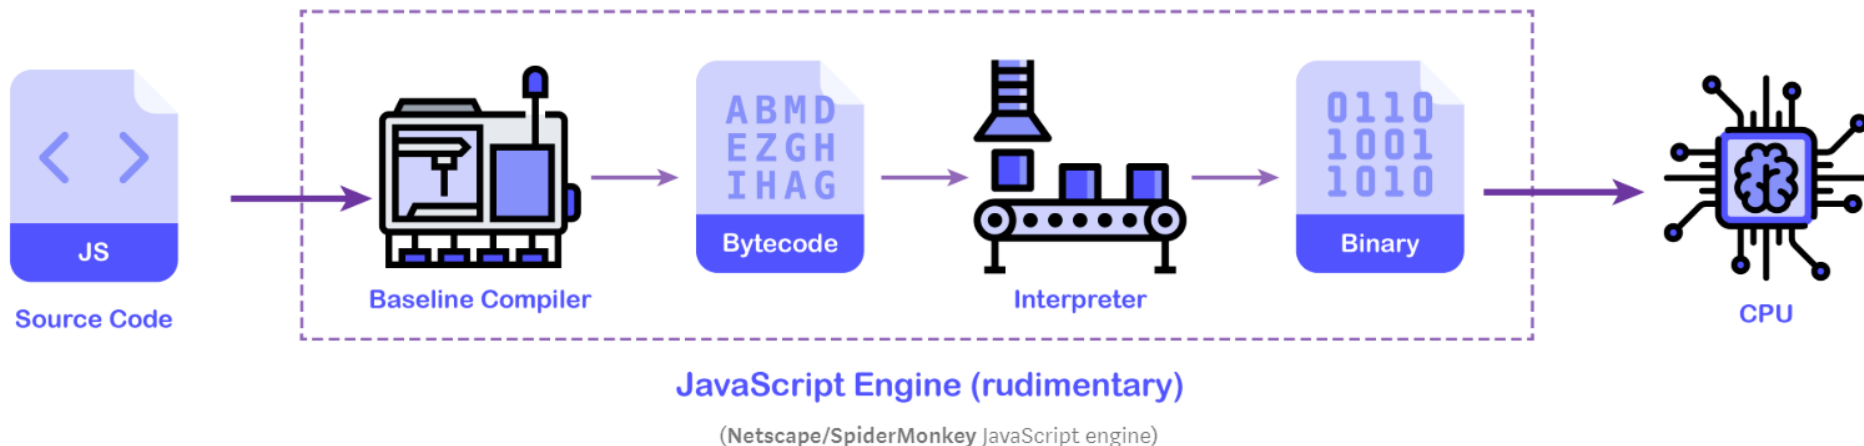
**LOW PERFORMANCE**

- 3D games? image/video editing?
- The cost of downloading, parsing, and compiling very large JavaScript applications can be prohibitive.
- Mobile and other resource-constrained platforms?

# How does JavaScript work?

- Every browser provides a JavaScript engine that runs the JavaScript code.

## *Rudimentary*

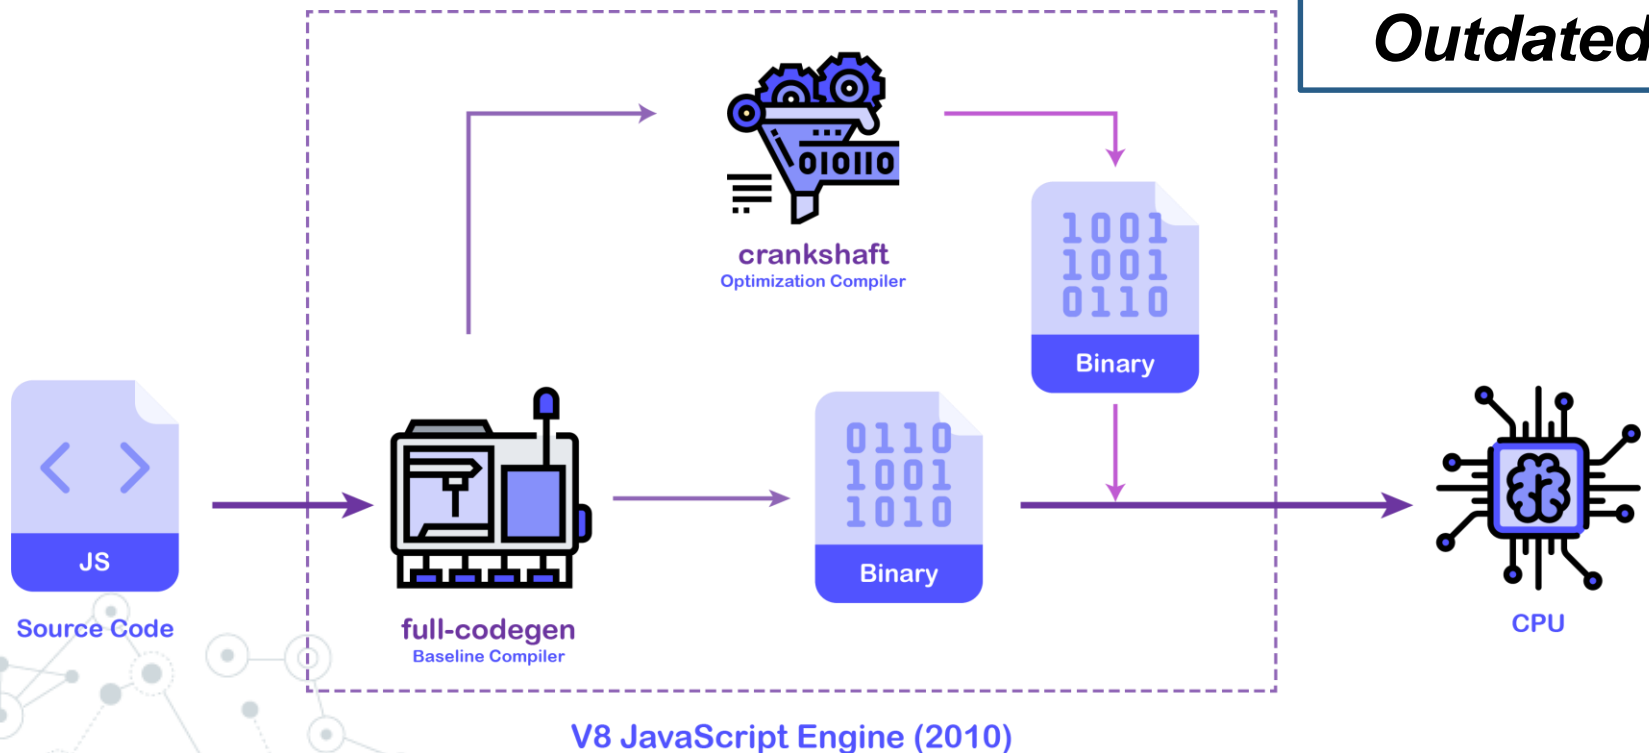


- When it comes to a highly dynamic and interactive web application, the user experience is very poor with this model.

# How does JavaScript work?

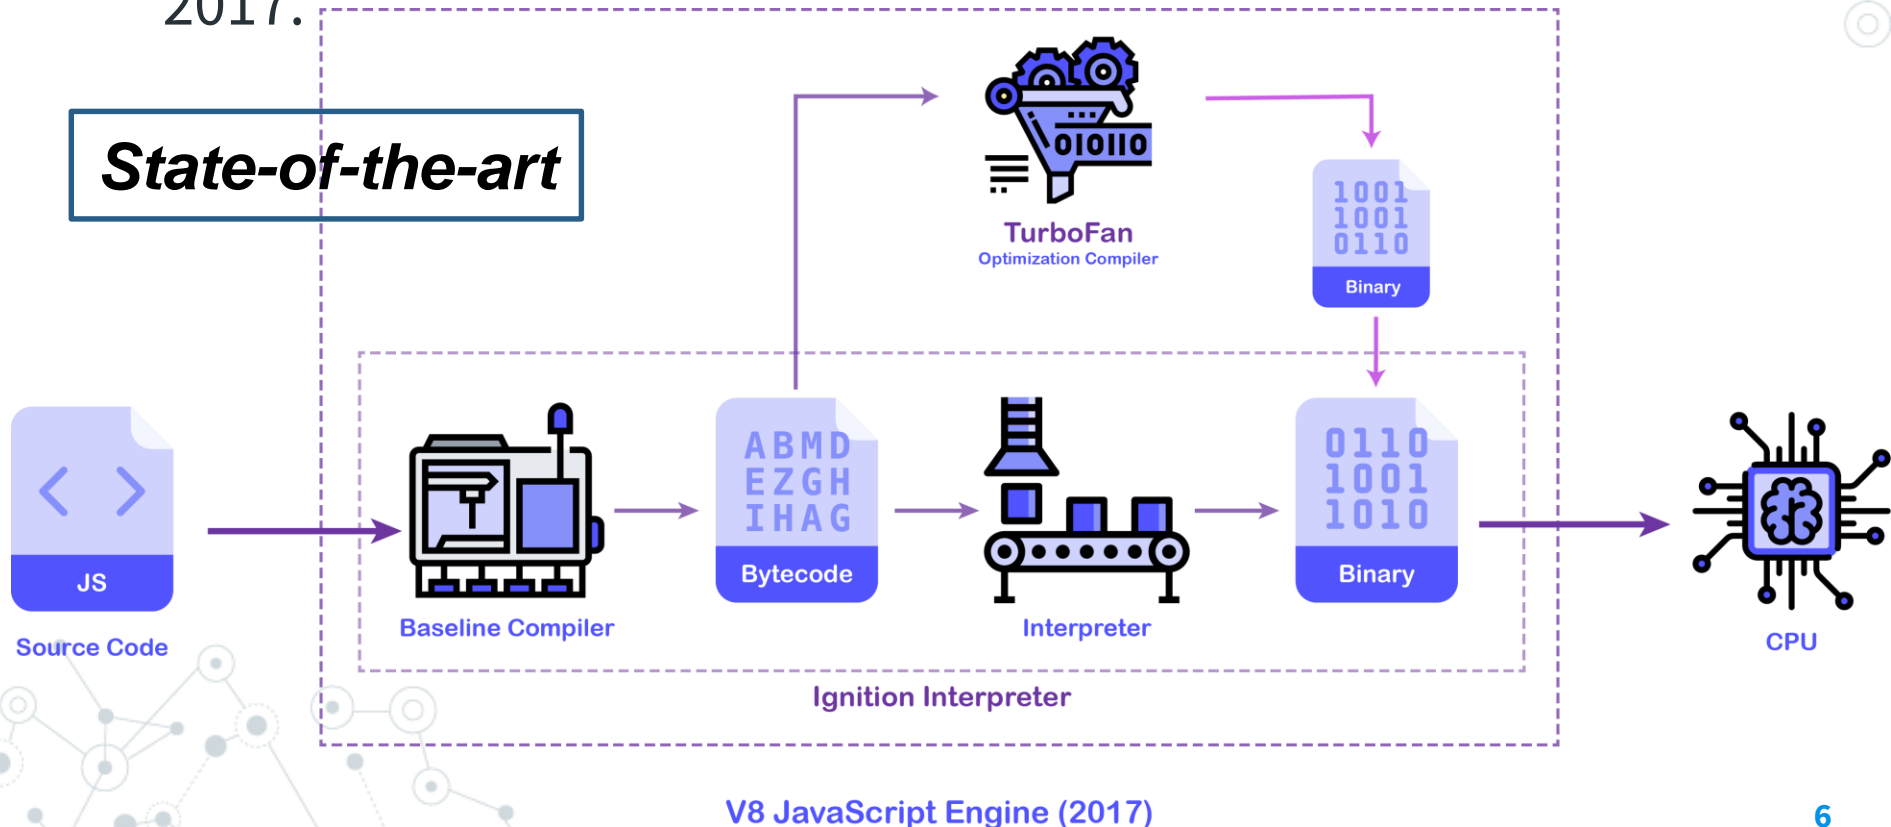
- This problem was faced by Google's Chrome browser while displaying Google Maps on the web.
- To increase the JavaScript performance on the web, they used the V8 JavaScript engine.


***Outdated***



# How JavaScript is optimized?

- The V8 team created a new version of the V8 engine from the ground up.
- This new version of the JavaScript engine was released in 2017.



A decorative background featuring a network diagram. It consists of numerous nodes, represented by circles of varying sizes and shades of gray, connected by thin, light gray lines. Some nodes are highlighted with a solid blue dot, and others are enclosed in a blue outline. The network is distributed across the slide, with a denser cluster on the left side and a more sparse arrangement on the right.

# WebAssembly: Game Changer!

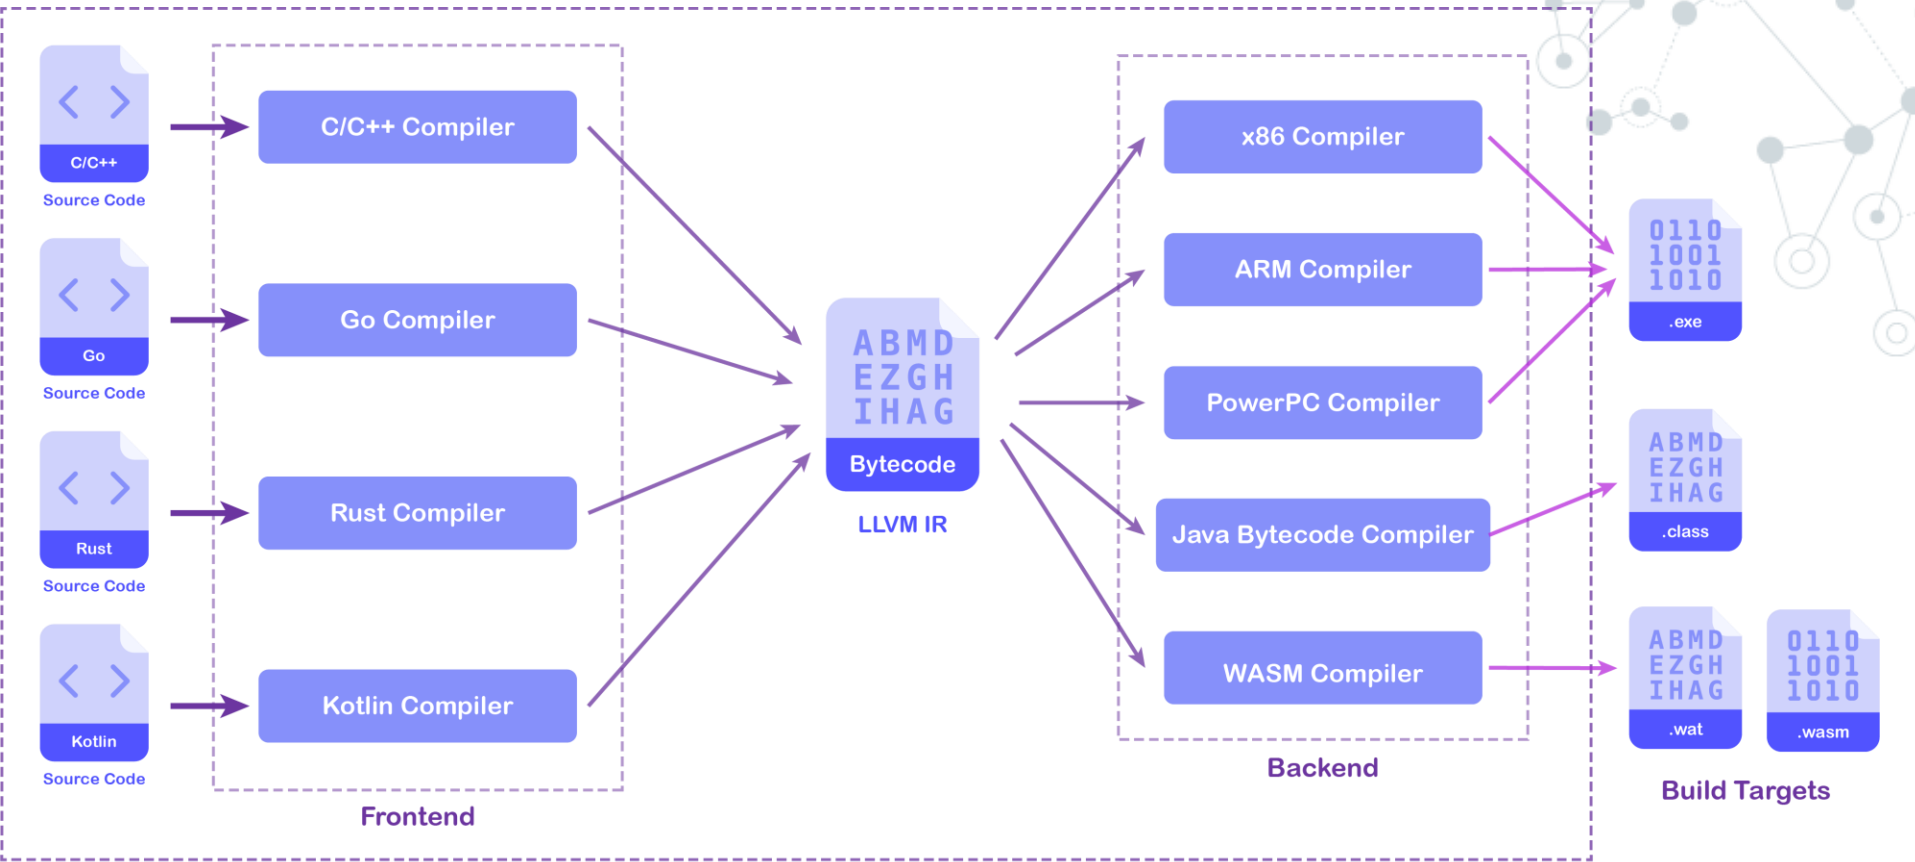
# WebAssembly: Game Changer!

- WebAssembly, a bytecode standard for web browsers.
- Announced in 17 June 2015.
- Developed by WebAssembly Working Group.
- Build target
- Binary format
- Supports JS.



# WebAssembly: Game Changer!

- A language for the web.
- Compiled from other languages.
- Offers maximized, reliable performance.
- Not a replacement for JS.
- It is meant to augment the things that JS was never designed to do.



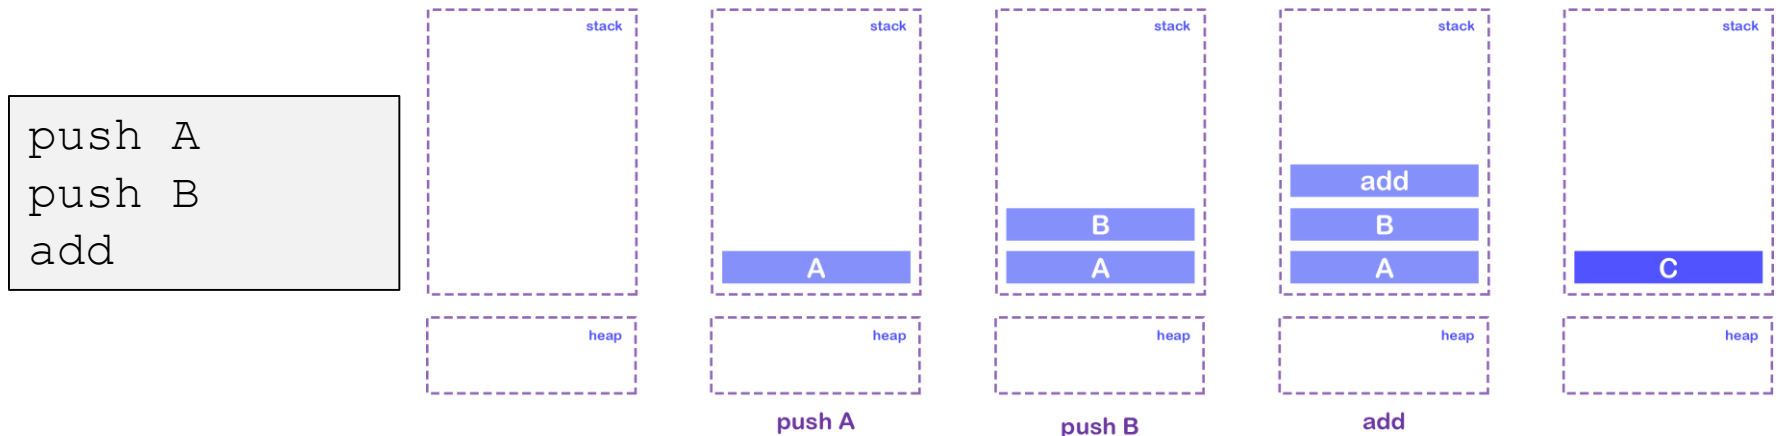
LLVM

# Main Advantages

- 1. Performance:** WebAssembly offers strong type guarantees, it gives more consistent and reliable performance than JS.
- 2. Portability:** because you can compile from other languages, you can bring open source libraries built in languages like C++.
- 3. Flexibility:** the ability to write in other languages.
  - To date, JS has been the only fully supported option.
  - Now with WebAssembly, you get more choice.

# WebAssembly is a Stack Machine

- Stack machine is a virtual machine (like a processor) that takes one instruction at a time and pushes it on the stack.
- When an operation is needed to be performed, it will pop values from the stack and compute the result.
- Example: compute the sum of two integers:



# WebAssembly is a Stack Machine

- Here's an example of a WebAssembly function that computes the quantity  **$F(x)=2x^2+1$** .

```
(func $f (param $x i32) (result i32)
    ;; stack: []
    (get_local $x) ;; stack: [x]
    (get_local $x) ;; stack: [x, x]
    (i32.mul)      ;; stack: [x*x]

    (i32.const 2)  ;; stack: [2, x*x]
    (i32.mul)      ;; stack: [2*x*x]

    (i32.const 1)  ;; stack: [1, 2*x*x]
    (i32.add))     ;; stack: [2*x*x+1]
```

# WebAssembly is a Stack Machine

- First, the syntax: this is the WebAssembly text format, which is derived from S-expressions.
- The same format commonly used for the Lisp programming.
- An S-expression is either
  - an “atom” (e.g. `get_local`, `2`, or `$x`),
  - or a list of S-expressions surrounded by parentheses (e.g. `(get_local $x)`, `(param $x i32)`).

# WebAssembly is a Stack Machine

- **MOST** instructions in WebAssembly modify the value stack in some way.
- Stack uses reverse polish notation

```
(i32.const 1)  
(i32.const 2)  
(i32.sub)
```

# F(x) in a register-based notation

```
(i3f:
    movl    %edi, -4(%rbp) ; *(rbp-4) = x
    shll    $1, %edi      ; edi = edi << 1 (= edi * 2)
    imull   -4(%rbp), %edi ; edi = edi * *(rbp-4)
    addl    $1, %edi      ; edi = edi + 1
    movl    %edi, %eax    ; eax = edi
    retq                               ; return eax2.const 1)
(i32.const 2)
(i32.sub)
```

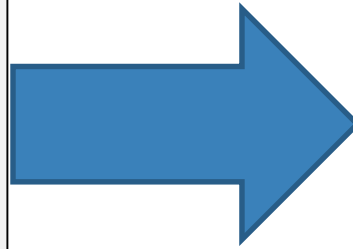
- Of course, x86 still has a stack (e.g. `-4(%rbp)`), but it uses that in tandem with registers.



# WebAssembly is Build Target

- Binary format - WASM

```
get_local 0
i64.eqz
if (result i64)
    i64.const 1
else
    get_local 0
    get_local 0
    i64.const 1
    i64.sub
    call 0
    i64.mul
end
```



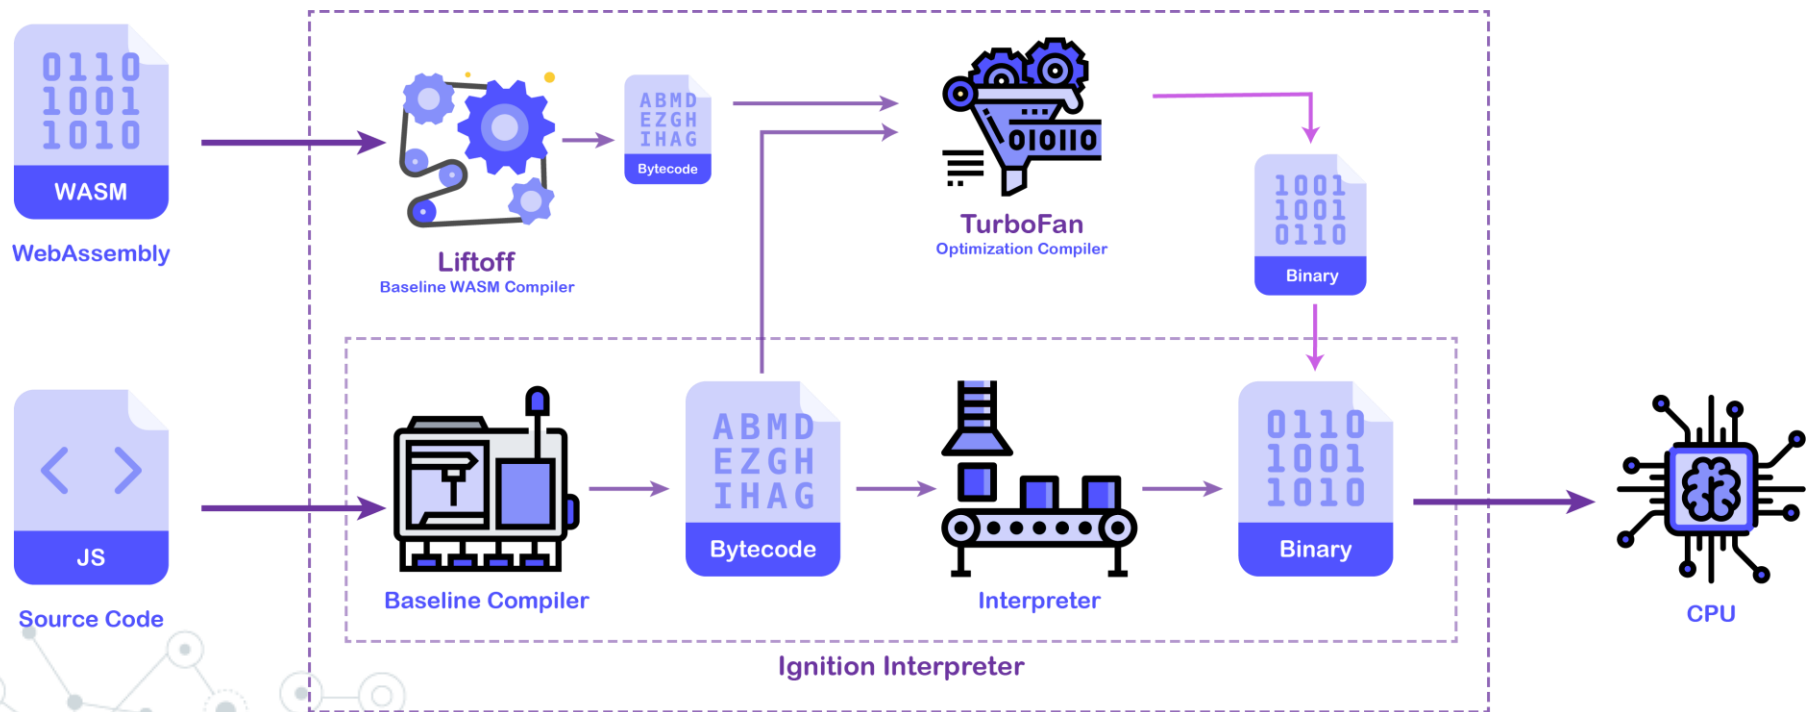
```
20 00
50
04 7E
42 01
05
20 00
20 00
42 01
7D
10 00
7E
0B
```

# How WebAssembly works?

- Every browser has a JavaScript engine which runs JavaScript.
- How it can run WebAssembly binary instructions?
- Browsers have introduced a new baseline compiler to compile WebAssembly to a bytecode that JavaScript interpreter can understand.

# How WebAssembly works?

- V8 integrated Liftoff, their WebAssembly baseline compiler whose job is to compile WebAssembly into an unoptimized bytecode as quick as possible.



# Example with WAT (WebAssembly text format)

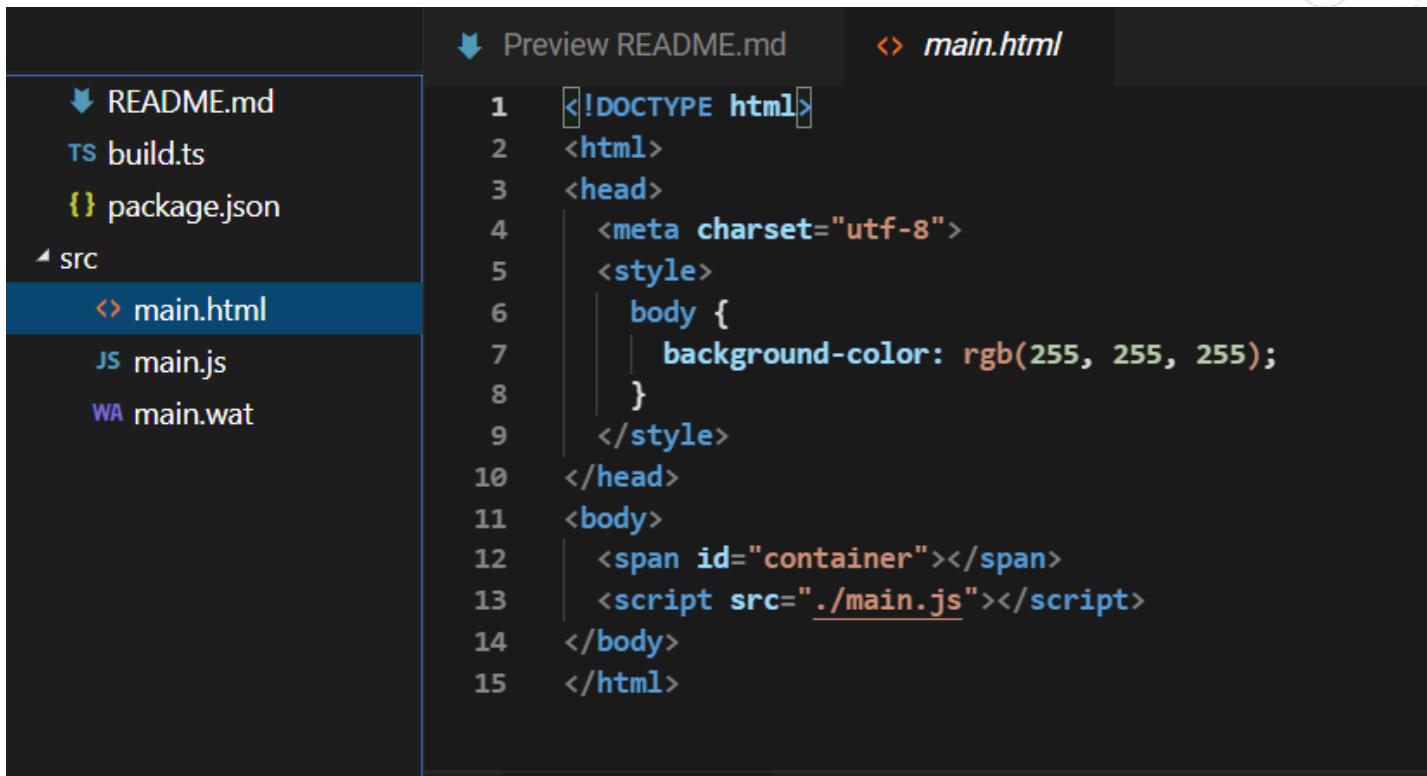
- Create an Empty Wat Project using

<https://webassembly.studio/>

- The project mainly contain three files



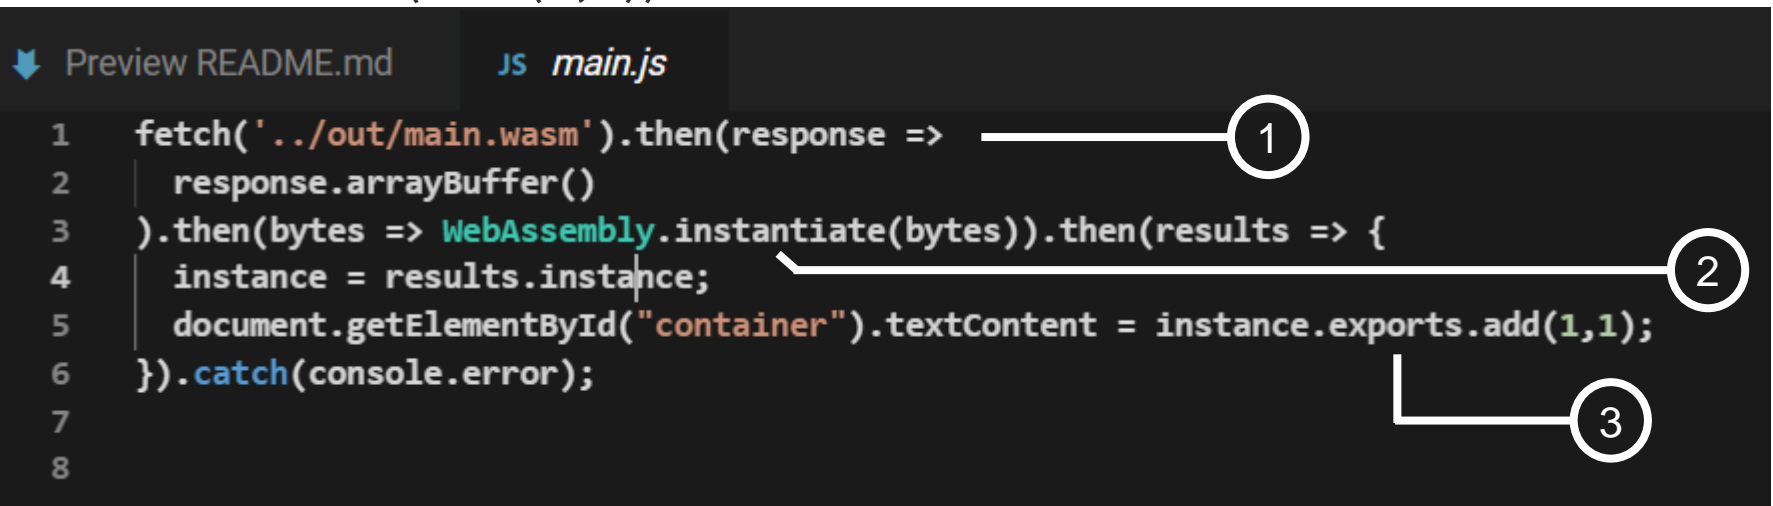
# HTML



```
1 <!DOCTYPE html>
2 <html>
3 <head>
4   <meta charset="utf-8">
5   <style>
6     body {
7       background-color: rgb(255, 255, 255);
8     }
9   </style>
10 </head>
11 <body>
12   <span id="container"></span>
13   <script src="./main.js"></script>
14 </body>
15 </html>
```

# JavaScript

1. To use a .wasm file in Javascript, .wasm file needs to be loaded as an external resource.
2. A loaded .wasm binary file needs to be compiled and linked (aka instantiated).
3. after creating an instance, the exported functions of the instance can be called through Javascript, as a normal JS functions (add(1,1))



```
Preview README.md JS main.js
1 fetch('../out/main.wasm').then(response => 1
2   response.arrayBuffer()
3 ).then(bytes => WebAssembly.instantiate(bytes)).then(results => {
4   instance = results.instance; 2
5   document.getElementById("container").textContent = instance.exports.add(1,1);
6 }).catch(console.error);
7
8
```

# WebAssembly (WAT)

- The module implements "add" function and exports it.

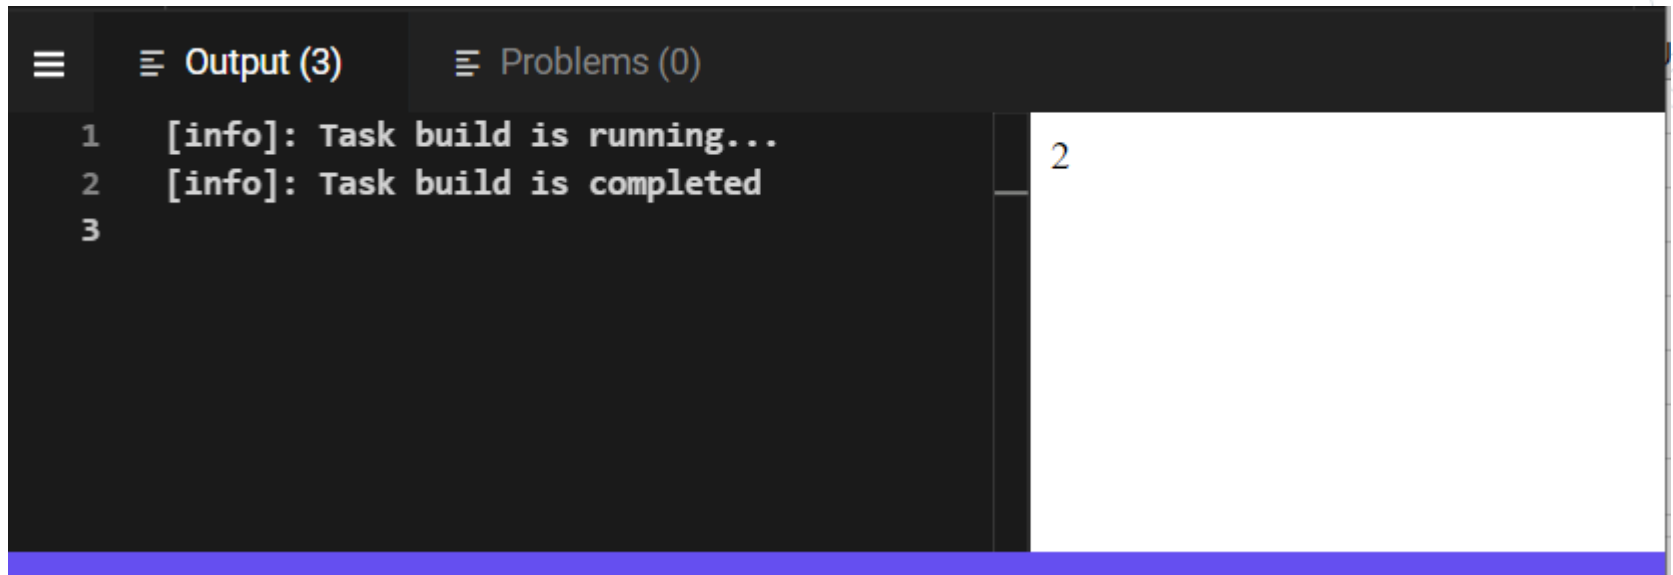
↓ Preview README.md

WA *main.wat*

```
1 (module
2   (func $add (param $lhs i32) (param $rhs i32) (result i32)
3     get_local $lhs
4     get_local $rhs
5     i32.add)
6   (export "add" (func $add)))
7 )
```

# Example with WAT

- Running the example should prints 2 to the browser



The screenshot shows a web browser's developer console with a dark theme. At the top, there are three tabs: a hamburger menu icon, 'Output (3)', and 'Problems (0)'. The 'Output (3)' tab is active, displaying a list of log messages. On the left side of the log, there are line numbers 1, 2, and 3. The log messages are: '[info]: Task build is running...' on line 1, '[info]: Task build is completed' on line 2, and an empty line on line 3. To the right of the log, there is a large white rectangular area representing the browser's output. The number '2' is printed in this area, corresponding to the completion message in the log.

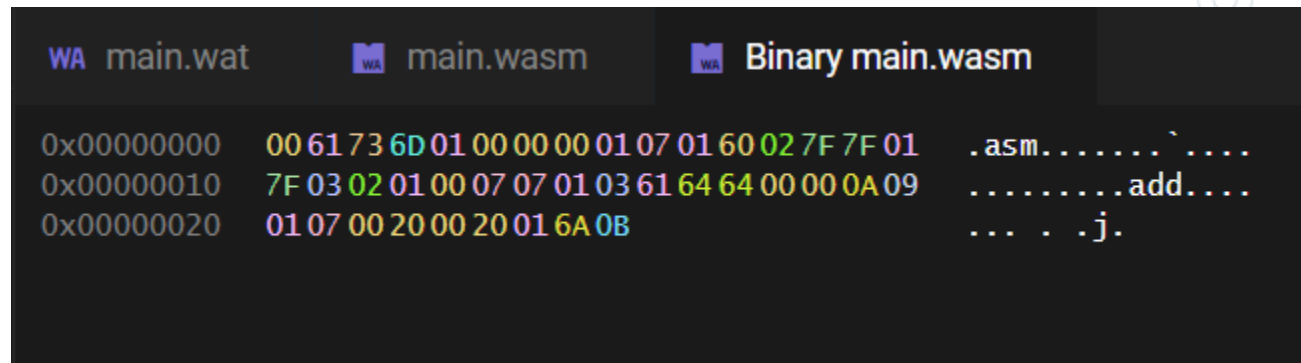
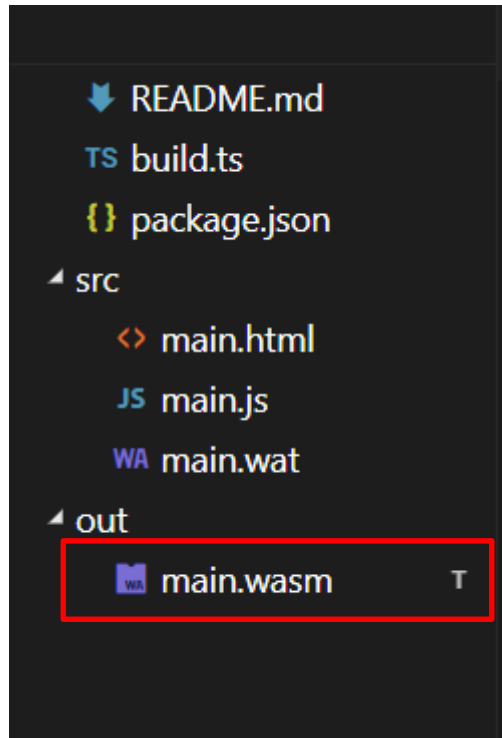
```
1 [info]: Task build is running...
2 [info]: Task build is completed
3
```

2



# Example with WAT

- Building the project creates main.wasm



# Another Example (Hello World) with WAT

0x00000000	00 61 73 6D 01 00 00 00	01 88 80 80 80 80 02 60	.asm.....`	(module
0x00000010	01 7F 00 60 00 00 02 8D	80 80 80 80 01 03 73 79	...`.....sy	(type \$type0 (func (param i32)))
0x00000020	73 05 70 72 69 6E 74 00	00 03 82 80 80 80 00 01	s.print.....	(type \$type1 (func))
0x00000030	01 05 86 80 80 80 00 01	01 01 C8 01 C8 01 07 91 80	.....	(import "sys" "print" (func \$import0 (param i32)))
0x00000040	80 80 00 02 06 6D 65 6D	6F 72 79 02 00 04 6D 61	.....memory...ma	(memory \$memory0 200 200)
0x00000050	69 6E 00 01 0A 8C 80 80	80 80 01 86 80 80 80 00	in.....	(export "memory" (memory \$memory0))
0x00000060	00 41 00 10 00 0B 0B 93	80 80 80 80 01 00 41 00	.A.....A.	(export "main" (func \$func1))
0x00000070	0B 0D 48 65 6C 6C 6F 2C	20 77 6F 72 6C 64 00	..Hello, world.	(func \$func1

```
(i32.const 0
  call $import0
)
(data (i32.const 0)
  "Hello, world\00"
)
)
```

# Can we use WebAssembly?

- Shipped in all major browsers.
- The first new language to ship in every major browser since JS.



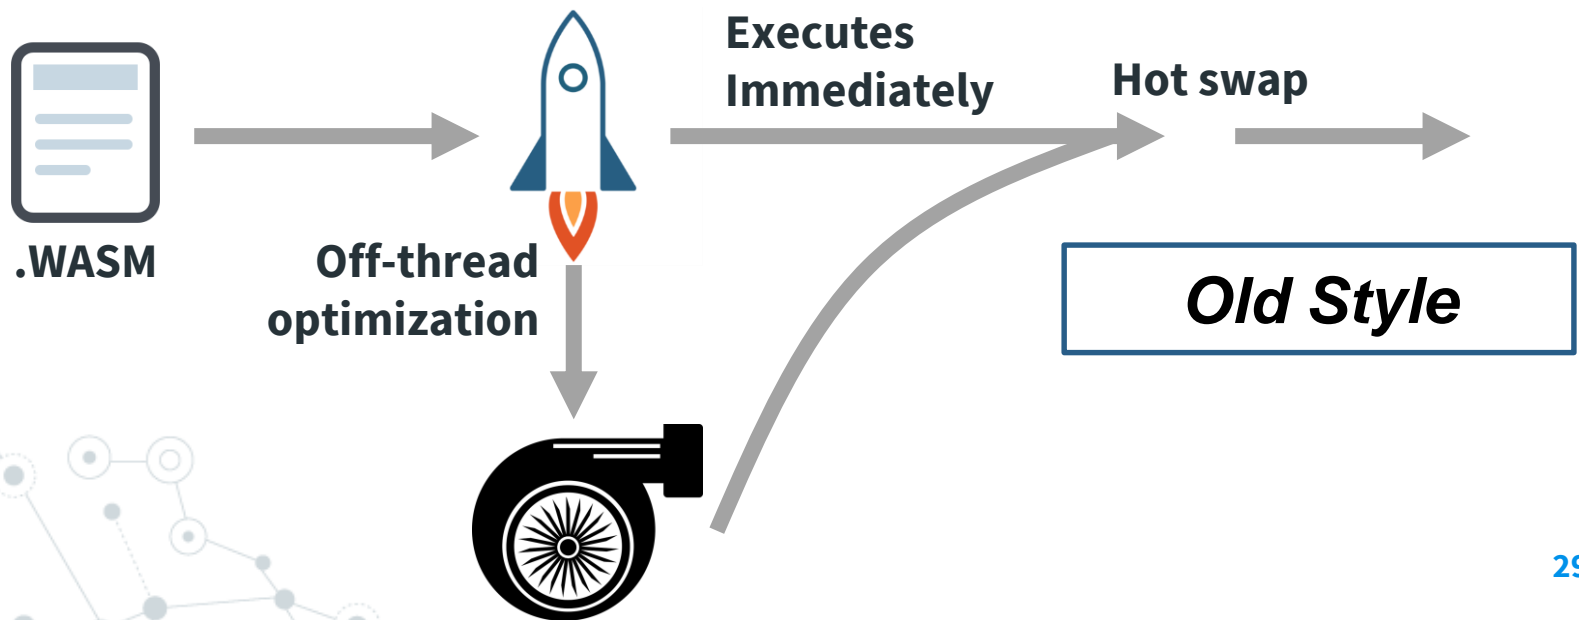
A decorative background featuring a network diagram. It consists of numerous nodes, represented by circles of varying sizes and colors (gray, blue, and white with blue outlines), connected by thin gray lines. The nodes are distributed across the slide, with a higher concentration in the top-left and bottom-right corners, creating a sense of connectivity and complexity.

# Why WebAssembly is Fast?

# Improvements to WASM

## 1- Implicit Caching (faster startup time):

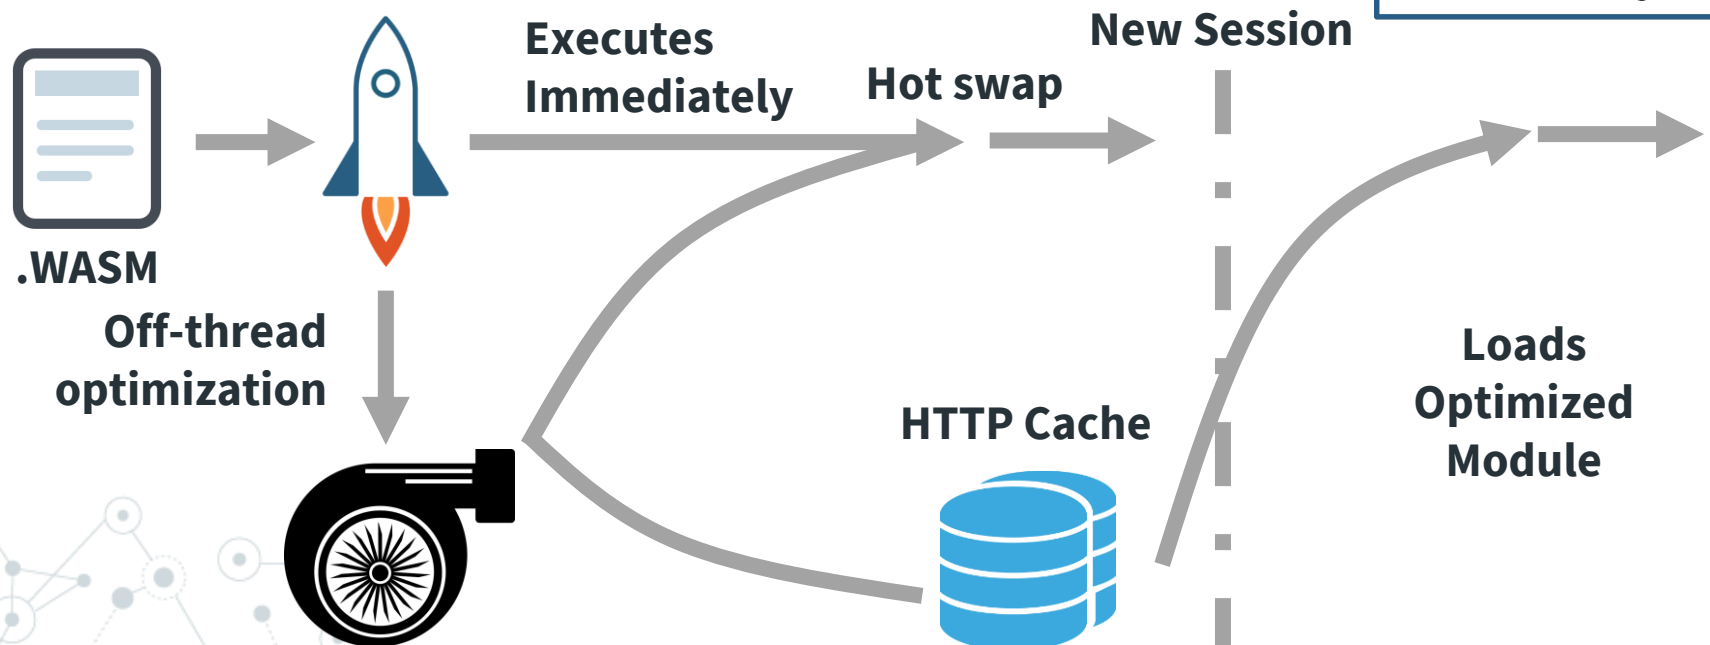
- When a site loads a WebAssembly module, it first goes into the Liftoff compiler to execute it.
- Then, the code is further optimized off the main thread through the turbo fan optimizing compiler.
- Then, the results are hot swapped in when ready.



# Improvements to WASM

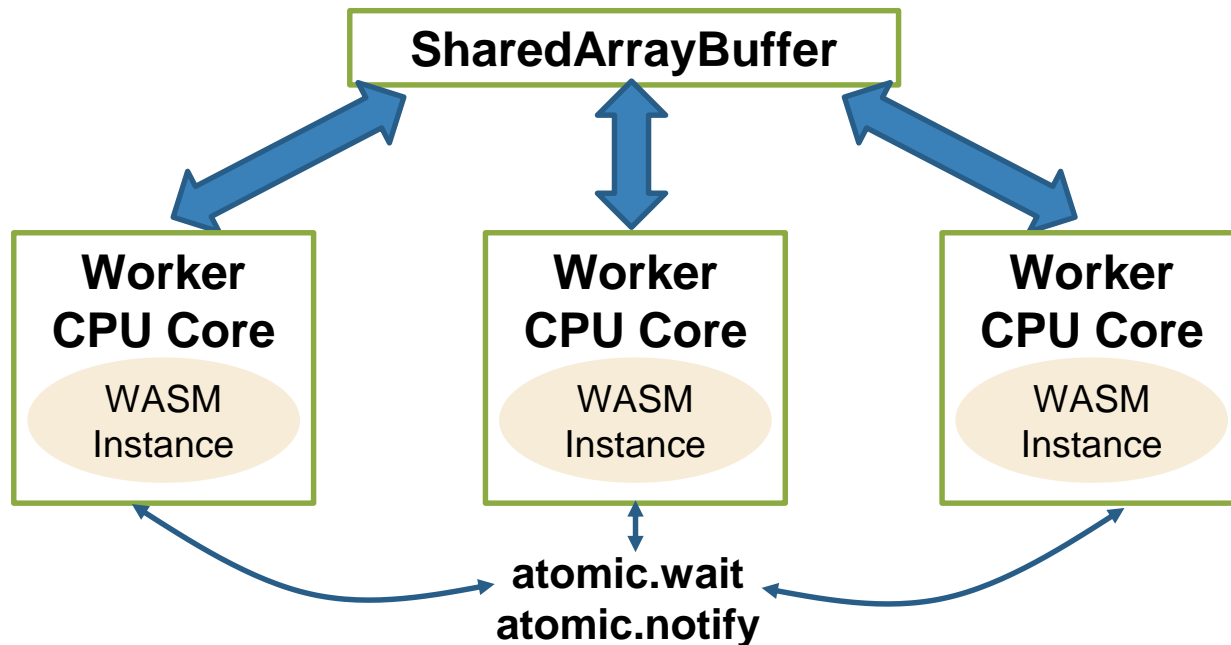
- Now, with implicit caching, we cache that optimized WebAssembly module directly in the HTTP cache.
- Then, after the user leaves the page and comes back, we load that optimized module directly from the cache, resulting in immediate top tier performance.

***New Style***



# Improvements to WASM

## 2- New Features: *WebAssembly Threads*



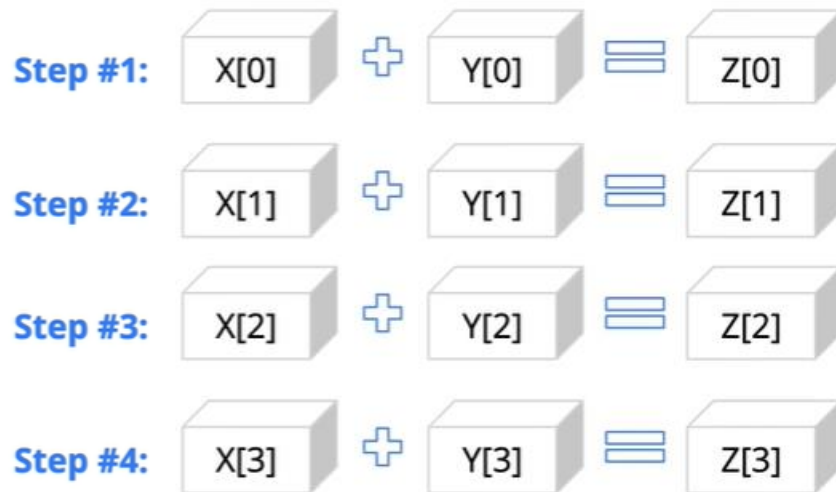
- WebWorkers: allows WebAssembly to run on different CPU cores.
- SharedArrayBuffer: allows WebAssembly to operate on the same piece of memory.
- Atomic Operations: synchronizing WebAssembly so that things happen in the right order.

# Improvements to WASM

## 2- New Features: *SIMD* (Single Instruction Multiple Data)

```
for (i = 0; i < 4; i++)  
    z[i] = x[i] + y[i]
```

### *Without SIMD*



### *With SIMD*

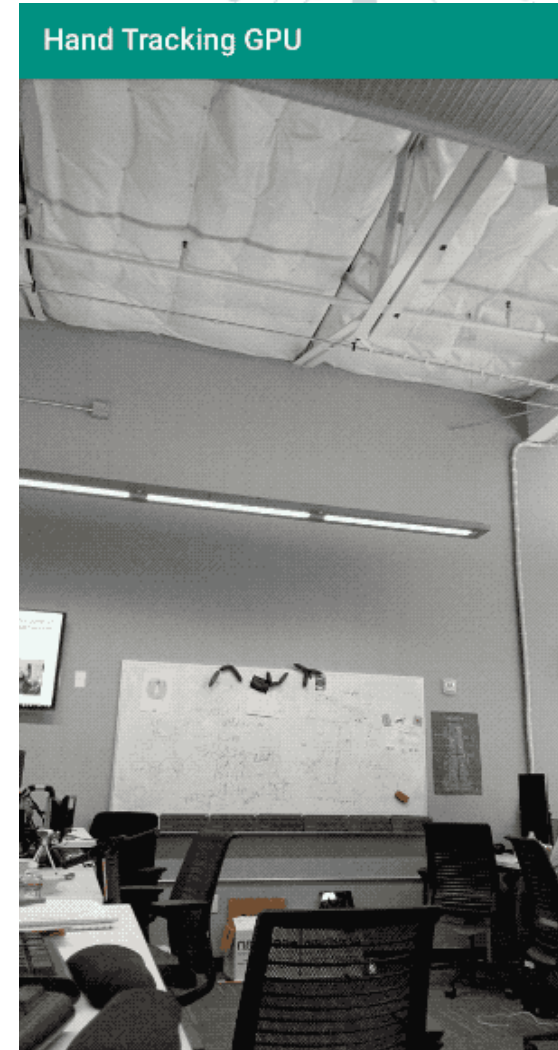




# Improvements to WASM

- The CPU is able to vectorize these elements and just take a single CPU operation to add them together.
- One example of SIMD is the hand tracking system:  
[Mediapipe.page.link/web](https://mediapipe.page.link/web)
- Without SIMD: we are getting about three frames per second
- With SIMD, we get a much smoother 15 frames per second.

***Please start the slide show to run the demo***



# Rust & WebAssembly Demo

- <https://m1el.github.io/wasm-asteroids/demo/demo.html>
- How will you make project 3 this exciting?