


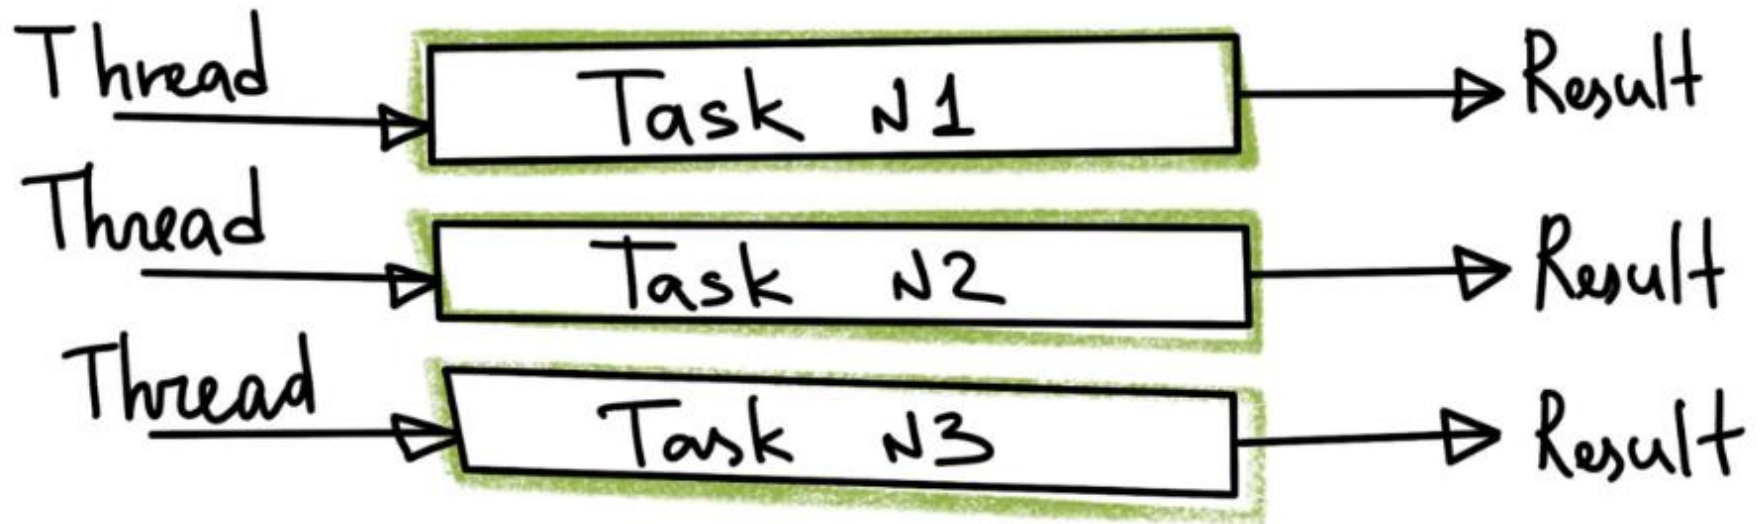
Lecture 16

The world is really Asynchronous!


Single-threaded



Multi-threaded



Asynchrony



Asynchrony and context switching

- Asynchronous programming was actually designed for an entirely different problem: CPU context switching.
- When you have multiple threads running, each CPU core can still only run one thread at a time.
- In order to allow all threads/processes to share resources, the CPU switches context very often.
- To oversimplify things, the CPU, at a random interval, saves all the context info of a thread and switches to another thread.
- The CPU is constantly switching between your threads in non-deterministic intervals.

Asynchrony and context switching

- Asynchronous programming is essentially cooperative multitasking with user-space threading.
- The application manages the threads and context switching rather than the CPU.

In an asynchronous world, context is switched at defined switch points.

Synchronous vs Asynchronous

- The asynchronous model performs best when:
 1. There are a large number of tasks.
 2. The tasks perform lots of I/O, causing the synchronous program to waste lots of time.
 3. The tasks are largely independent of one another so there is little need for inter-task communication.

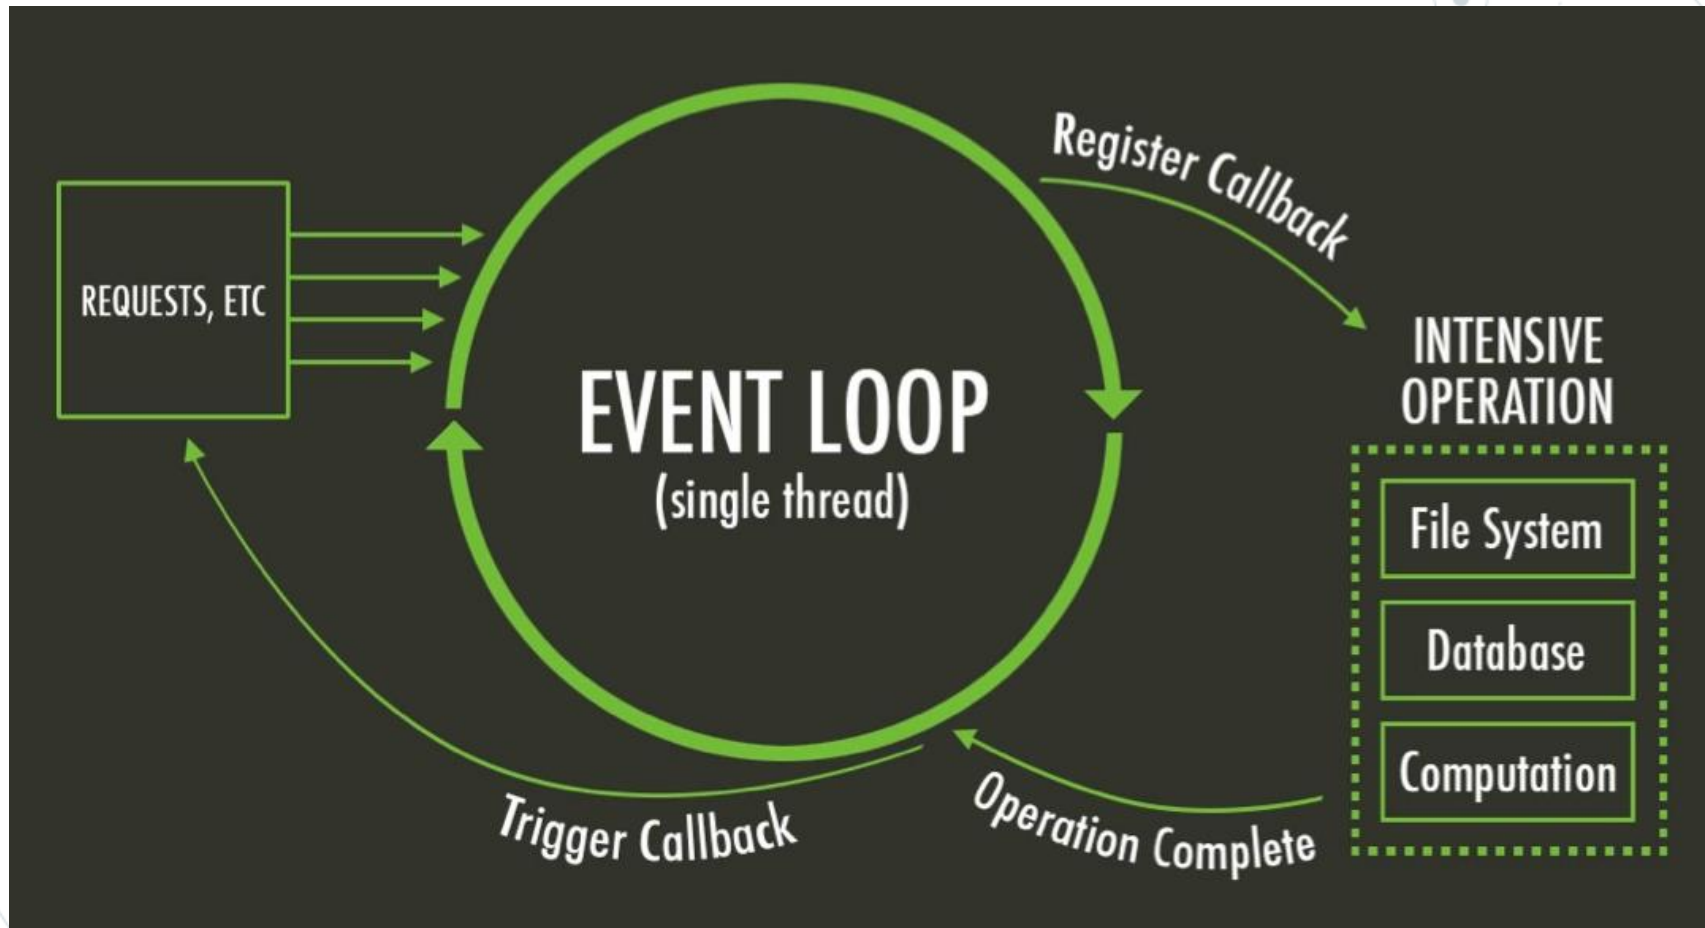
Synchronous vs Asynchronous

- These conditions almost perfectly characterize a typical busy server (like a webserver) in a client-server environment.
- Each task represents one client request with I/O in the form of receiving the request and sending the reply.

Synchronous vs Asynchronous

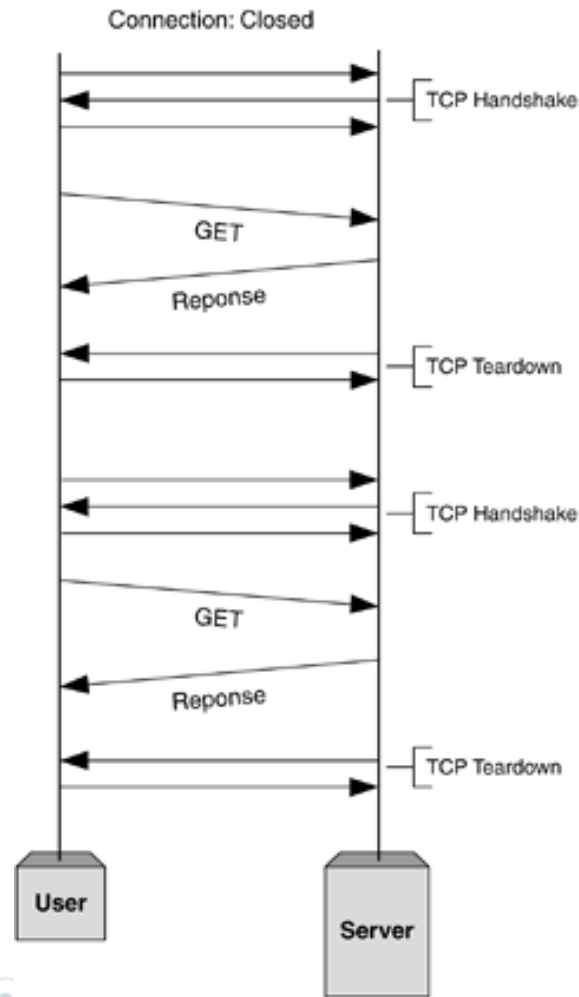
- However, as the number of threads increases, your server may start to experience performance problems.
- With each new thread, there is some memory overhead associated with the creation and maintenance of the thread.
- Another performance gain from the asynchronous model is that it avoids context switching.
- Every time the OS transfers control over from one thread to another it has to save all the relevant registers, memory map, stack pointers, CPU context etc.
- The overhead of doing this can be quite significant.

Event Loop (Remember GUI)

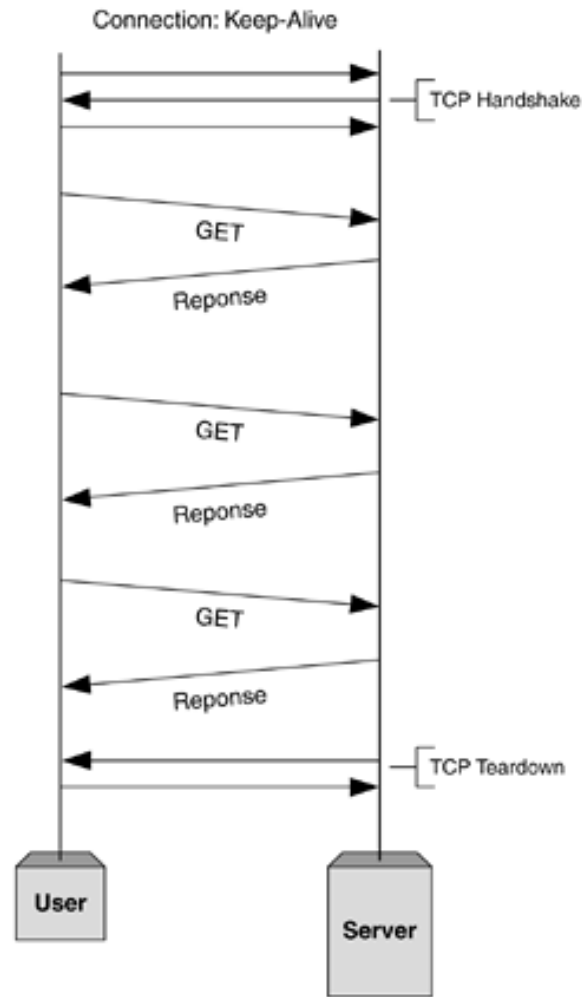


Networking, HTTP

HTTP 1 and 1.1





In a default HTTP/1.0 session, the TCP connection will be torn down and re-established between each HTTP GET request.



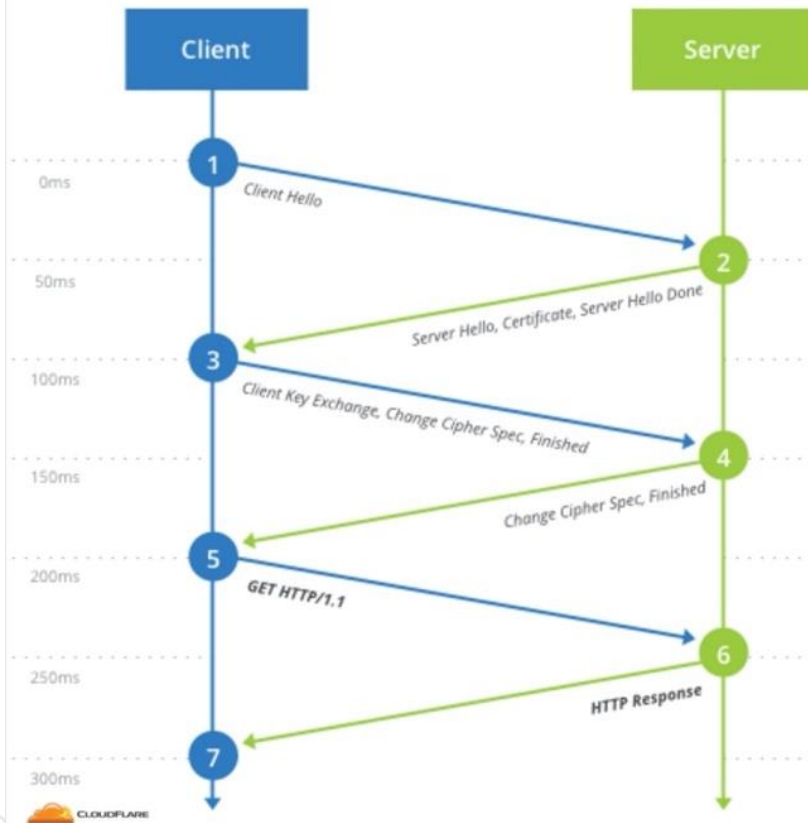
In a default HTTP/1.1 session, a single TCP connection will be held, open and multiple GET requests will be passed across.

HTTP 1.1 and 2.0

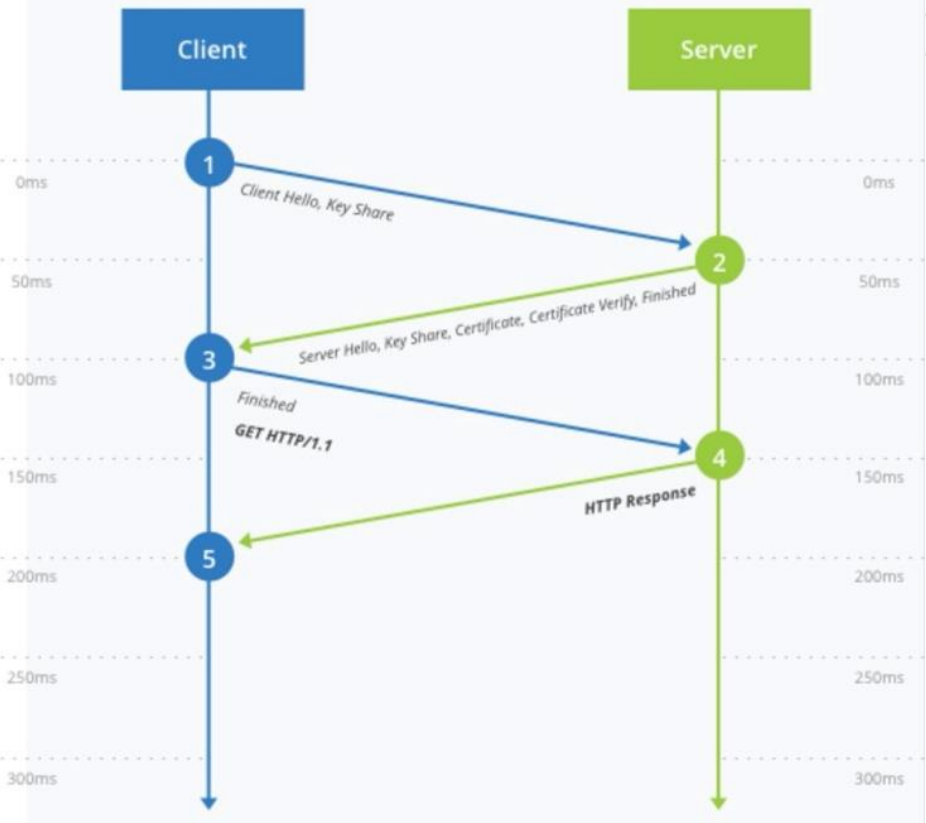


TLS 1.2 and 1.3

TLS 1.2 (Full Handshake)



TLS 1.3 (Full Handshake)



HTTPS 2.0 and 3.0

