

Lecture 19

Async I/O with Tokio

What is Tokio?

- Tokio is an event-driven, non-blocking I/O platform for writing asynchronous applications with RUST.
- The `async-std` library provides similar functionality.
- Network services are the most common domain for a non-blocking I/O system.
- Tokio is the most popular and established of these systems today.

An asynchronous hello world application:

- `#[tokio::main]` adds boilerplate code, now looks like std calls

```
use tokio::io;

#[tokio::main]
async fn main() -> Result<(), std::io::Error> {
    let mut stdout = io::stdout();
    let mut hello: &[u8] = b"Hello, world!\n"; //b? a slice
    io::copy(&mut hello, &mut stdout).await?;
    Ok(())
}
```

- We are generating a **Future**, so our main function is async.
- Since main is async, we need to use an executor to run it. That's why we use the `#[tokio::main]` attribute.
- Since performing I/O can fail (?), we return a Result.

.await?

- .await?:
 - .await for chaining together Futures.
 - ? for error handling.
- We use `tokio::io::stdout()` to get access to some value that lets us interact with standard output.
- looks really similar to `std::io::stdout()`.
- a large part of the tokio API is simply async-ifying things from std.

tokio::io::copy

- An asynchronous version of `std::io::copy`.
- Instead of working with the *Read* and *Write* traits, this works with their async cousins: *AsyncRead* and *AsyncWrite*.
- A byte slice (`&[u8]`) is a valid *AsyncRead*, so we're able to store our input there.
- `Stdout` is an *AsyncWrite*.

Example 1

- Modify this application so it copies the entire contents of standard input to standard output.

```
use tokio::io;

#[tokio::main]
async fn main() -> Result<(), std::io::Error> {
    let mut stdin = io::stdin();
    let mut stdout = io::stdout();
    io::copy(&mut stdin, &mut stdout).await?;
    Ok(())
}
```

Spawning Processes

- Tokio provides a `tokio::process` module which resembles the `std::process` module.

```
use tokio::process::Command;

#[tokio::main]
async fn main() -> Result<(), std::io::Error> {
    Command::new("echo").arg("Hello, world!").spawn()?.await?;
    Ok(())
}
```

Spawning Processes

```
Command::new("echo").arg("Hello, world!").spawn()?.await?;
```



1. Create a new Command to run echo
2. Give it the argument "Hello, world!"
3. Spawn this, which may fail
4. ?: if it fails, return the error. Otherwise, return a Future
5. Using the .await: wait until that Future completes, and capture its Result
6. ?: if that Result is Err, return that error.

Delay_for

```
use tokio::time;
use tokio::process::Command;
use std::time::Duration;

#[tokio::main]
async fn main() -> Result<(), std::io::Error> {
    Command::new("date").spawn()?.await?;
    time::delay_for(Duration::from_secs(1)).await;
    Command::new("date").spawn()?.await?;
    time::delay_for(Duration::from_secs(1)).await;
    Command::new("date").spawn()?.await?;
    Ok(())
}
```

tokio::time::interval

- Create a stream of “ticks” .
- This program will keep calling date once per second until it is killed:

```
use tokio::time;
use tokio::process::Command;
use std::time::Duration;

#[tokio::main]
async fn main() -> Result<(), std::io::Error> {
    let mut interval = time::interval(Duration::from_secs(1));
    loop {
        interval.tick().await;
        Command::new("date").spawn()?.await?;
    }
}
```

Time to spawn

```
use std::time::Duration;
use tokio::process::Command;
use tokio::task;
use tokio::time;

#[tokio::main]
async fn main() -> Result<(), std::io::Error> {
    task::spawn(dating()).await??;
    Ok(())
}

async fn dating() -> Result<(), std::io::Error> {
    let mut interval = time::interval(Duration::from_secs(1));
    loop {
        interval.tick().await;
        Command::new("date").spawn()?.await?;
    }
}
```

??

- Calling spawn gives us back a `JoinHandle<T::Output>`.

```
pub fn spawn<T>(task: T) -> JoinHandle<T::Output>;  
impl<T> Future for JoinHandle<T> {  
    type Output = Result<T, JoinError>;  
}
```

- The Future we provide as input is `dating()`, which has an output of type `Result<(), std::io::Error>`.
- The type of `task::spawn(dating())` is `JoinHandle<Result<(), std::io::Error>>`.
- `JoinHandle` implements `Future`, so we end up with `Result<Result<(), std::io::Error>, JoinError>`.
- The first `?` deals with the outer `Result`.
- The second `?` deals with the `std::io::Error`.

Synchronous code

- Tokio's performs blocking calls with the `spawn_blocking` function.
- The `task::spawn_blocking` function is similar to the `task::spawn` function
- But rather than spawning an non-blocking future on the Tokio runtime, it instead spawns a blocking function on a dedicated thread pool for blocking tasks.

Let's network!

```
use tokio::io;
use tokio::net::{TcpListener, TcpStream};
```

```
#[tokio::main]
async fn main() -> io::Result<()> {
    let mut listener =
    TcpListener::bind("127.0.0.1:8080").await?;
```

TcpListener binds a socket. The binding itself is asynchronous, so we use `.await` to wait for the listening socket to be available.

```
    loop {        //we loop forever
        let (socket, _) = listener.accept().await?;
        echo(socket).await?;
    }
}
```

Accept new connections, and call our echo function and `.await` it.

```
async fn echo(socket: TcpStream) -> io::Result<()> {
    let (mut recv, mut send) = io::split(socket);
    io::copy(&mut recv, &mut send).await?;
    Ok(())
}
```

We use `tokio::io::split` to split up our `TcpStream` into read and write halves, and then pass those into `tokio::io::copy`.

Let's network!

- What should the behavior be if a second connection comes in while the first connection is still active?
- The program has just one task: awaits on each call to echo.
- So the second connection won't be serviced until the first one closes.
- *Modify the program above so that it handles concurrent connections correctly.*

Let's network!

- Solution: Wrap the echo call with `tokio::spawn`:

```
loop {  
    let (socket, _) = listener.accept().await?;  
    tokio::spawn(echo(socket));  
}
```


TCP client and ownership

- It will establish a connection to a hard-coded server.
- Copy all of stdin to the server, and then copy all data from the server to stdout.

```
use tokio::io;
use tokio::net::TcpStream;

#[tokio::main]
async fn main() -> io::Result<()> {
    let stream = TcpStream::connect("127.0.0.1:8080").await?;
    let (mut recv, mut send) = io::split(stream);
    let mut stdin = io::stdin();
    let mut stdout = io::stdout();

    io::copy(&mut stdin, &mut send).await?;
    io::copy(&mut recv, &mut stdout).await?;

    Ok(())
}
```

TCP client and ownership

- But it's limited.
- It only handles half-duplex protocols like HTTP, and doesn't actually support keep-alive in any way.
- Can we use spawn to run the two copys in different tasks?

```
let send = spawn(io::copy(&mut stdin, &mut send));  
let recv = spawn(io::copy(&mut recv, &mut stdout));  
  
send.await?;  
recv.await?;
```

- Note: In a half-duplex system, only one device can transmit at a time.

TCP client and ownership

- However, this doesn't compile:

```
error[E0597]: `stdin` does not live long enough
  --> src/main.rs:12:31
   |
12 |         let send = spawn(io::copy(&mut stdin, &mut send));
   |                                   ^^^^^^^^^^^^^^^^^
   |                                   |
   |                                   borrowed value does not
live long enough
   |                                   argument requires that `stdin` is
   |                                   borrowed for `'static`
19 |     }
   |     - `stdin` dropped here while still borrowed
```

- Copy Future does not own the stdin value.
- To fix this, we need to convince the compiler to make a Future that owns stdin.

TCP client and ownership

```
use tokio::io;
use tokio::spawn;
use tokio::net::TcpStream;

#[tokio::main]
async fn main() -> Result<(), std::io::Error> {
    let stream = TcpStream::connect("127.0.0.1:8080").await?;
    let (mut recv, mut send) = io::split(stream);
    let mut stdin = io::stdin();
    let mut stdout = io::stdout();

    let send = spawn(async move {
        io::copy(&mut stdin, &mut send).await
    });
    let recv = spawn(async move {
        io::copy(&mut recv, &mut stdout).await
    });

    send.await??;
    recv.await??;

    Ok(())
}
```

Playing with lines

- Let's build an async program that counts the number of lines on standard input.

```
use tokio::prelude::*;
use tokio::io::AsyncBufReadExt;

#[tokio::main]
async fn main() -> Result<(), std::io::Error> {
    let stdin = io::stdin();
    let stdin2 = io::BufReader::new(stdin);
    let mut count = 0u32;
    let mut lines = stdin2.lines();
    while let Some(_) = lines.next_line().await? {
        count += 1;
    }
    println!("Lines on stdin: {}", count);
    Ok(())
}
```

Playing with lines

- Let's take a list of file names as command line arguments, and count up the total number of lines in all the files.

```
use tokio::prelude::*;
use tokio::io::AsyncBufReadExt;

#[tokio::main]
async fn main() -> Result<(), std::io::Error> {
    let mut args = std::env::args();
    let _me = args.next(); // ignore command name
    let mut count = 0u32;
    for filename in args {
        let file = tokio::fs::File::open(filename).await?;
        let file2 = io::BufReader::new(file);
        let mut lines = file2.lines();
        while let Some(_) = lines.next_line().await? {
            count += 1;
        }
    }
    println!("Total lines: {}", count);
    Ok(())
}
```

Make it asynchronous

```
use tokio::prelude::*;
use tokio::io::AsyncBufReadExt;

#[tokio::main]
async fn main() -> Result<(), std::io::Error> {
    let mut args = std::env::args();
    let _me = args.next(); // ignore command name
    let mut tasks = vec![];
    for filename in args {
        tasks.push(tokio::spawn(async {
            let file = tokio::fs::File::open(filename).await?;
            let file2 = io::BufReader::new(file);
            let mut lines = file2.lines();
            let mut count = 0u32;
            while let Some(_) = lines.next_line().await? {
                count += 1;
            }
            Ok(count) as Result<u32, std::io::Error>
        }));
    }
    let mut count = 0;
    for task in tasks {
        count += task.await??;
    }
    println!("Total lines: {}", count);
    Ok(())
}
```

Playing with lines

- Let's change how we handle the count.
- Instead of `.awaiting` the count in the second for loop, let's have each individual task update a shared mutable variable.
- You should use an `Arc<Mutex<u32>>`.


```

#[tokio::main]
async fn main() -> Result<(), std::io::Error> {
    let mut args = std::env::args();
    let _me = args.next(); // ignore command name
    let mut tasks = vec![];
    let count = Arc::new(Mutex::new(0u32));
    for filename in args {
        let count = count.clone();
        tasks.push(tokio::spawn(async move {
            let file = tokio::fs::File::open(filename).await?;
            let file = io::BufReader::new(file);
            let mut lines = file.lines();
            let mut local_count = 0u32;
            while let Some(_) = lines.next_line().await? {
                local_count += 1;
            }
            let mut count = count.lock().await;
            *count += local_count;
            Ok(()) as Result<(), std::io::Error> }));
    }
    for task in tasks {
        task.await??;
    }
    let count = count.lock().await;
    println!("Total lines: {}", *count);
    Ok(())
}

```